

Quantum Programming Languages and Semantics

5 QWhile Language

This section introduces the core programming language we will use for quantum program logics. The language follows the *classical control + quantum data* paradigm: the program's control flow (sequencing, branching, looping, and the program counter itself) is entirely classical, while the data manipulated by commands are quantum states (density operators) on a fixed composite Hilbert space. Concretely:

Control is classical. The program counter is never put into superposition. The only source of probabilistic branching is *measurement*: a measurement produces a classical outcome, and the next command is chosen based on that outcome.

Data is quantum. At every point, the program carries a (possibly subnormalized) density operator ρ on a fixed global Hilbert space H . Primitive commands act *locally* on specified subsystems, leaving the rest untouched (up to the standard identity extension).

Registers and subsystem structure. Fix a finite set of register labels Reg and a family of finite-dimensional Hilbert spaces $\{H_x\}_{x \in \text{Reg}}$. The *global* state space is the tensor product

$$H \cong \bigotimes_{x \in \text{Reg}} H_x.$$

A *subsystem* is a subset of labels $s \subseteq \text{Reg}$. We write

$$H_s \cong \bigotimes_{x \in s} H_x, \quad \bar{s} := \text{Reg} \setminus s, \quad H \cong H_s \otimes H_{\bar{s}}.$$

The purpose is to make precise the standard programming intuition that a program acts on *named registers* and typically touches only *finitely many* of them. A concrete quantum program is written in terms of a finite collection of variables/registers/qubits, such as q_1, q_2, \dots, q_n , and every primitive command (unitary gates, measurements, initialization) targets a specified subset of these registers. Thus it is natural to model the program's data state as a single density operator on a tensor product space that factors according to these register boundaries. Moreover, working with explicit subsystems is essential for two semantic reasons:

Locality of commands. A command that is declared to act on subsystem s should leave the rest of the machine state unchanged; this is expressed later by cylindrical extension ($A^{(s)} = A \otimes \mathbf{1}_{H_{\bar{s}}}$).

Discarding and re-initialization. Initialization/reset operations need to “forget” whatever was previously stored in a register (including any entanglement with other registers). This is naturally described using the decomposition $H \cong H_s \otimes H_{\bar{s}}$ together with partial trace.

We tacitly fix a tensor-factor ordering/parenthesization convention so that expressions such as $H \cong H_s \otimes H_{\bar{s}}$ are well-defined (up to canonical unitary isomorphism). This lets us speak cleanly about commands that act only on s while leaving \bar{s} unchanged.

Program states (subnormalized density operators). To represent probabilistic branching without constantly renormalizing, we use *subnormalized* (partial) density operators:

$$\mathcal{D}^-(H) := \{\rho \in L(H) : \rho \succeq 0 \wedge \text{tr}(\rho) \leq 1\}.$$

Intuitively, $\text{tr}(\rho)$ is the *probability mass* of reaching ρ along a particular execution path. Thus the same symbolic ρ can simultaneously encode the post-state *and* the probability of being in that post-state. The zero operator 0 is allowed and represents an impossible (unreachable) state.

Local operators and cylindrical extension. A key design principle is *locality*: commands target a subsystem $s \subseteq \text{Reg}$ and act as the identity on the complement \bar{s} . If $A \in L(H_s)$ is a linear operator acting on subsystem s , its *cylindrical extension* (identity extension) to the global space is

$$A^{(s)} := A \otimes \mathbf{1}_{H_{\bar{s}}} \quad (\text{under the fixed identification } H \cong H_s \otimes H_{\bar{s}}).$$

Likewise, if U_s is unitary on H_s , then $U_s^{(s)}$ is unitary on H and represents “apply U_s on subsystem s and do nothing elsewhere.”

We also need a primitive way to *discard* part of a composite system. For $\rho \in L(H)$ and $s \subseteq \text{Reg}$, write $\text{tr}_s(\rho)$ for the partial trace over H_s . Then $\text{tr}_s(\rho) \in L(H_{\bar{s}})$ and preserves subnormalization:

$$\text{tr}(\text{tr}_s(\rho)) = \text{tr}(\rho).$$

Operationally, tr_s means “throw away subsystem s and keep only the reduced state of the rest.” A key operation in *qwhile* is to *discard* the old content of subsystem s and replace it by a fresh local state ρ_s on H_s . Intuitively, we first “throw away” subsystem s by taking the partial trace $\text{tr}_s(\rho)$, and then prepare a new state ρ_s on s . As a result, the register s is set to ρ_s regardless of its previous contents, and any entanglement between s and \bar{s} is removed. This is the subsystem-level generalization of the familiar qubit reset command $q := |0\rangle$.

5.1 QWhile Language

We now define a core language, *qwhile*, that is expressive enough to support quantum Hoare-style correctness reasoning and under-approximate incorrectness reasoning. The distinctive features are: (i) commands act on designated subsystems; (ii) branching/looping is controlled by measurement outcomes; (iii) the statement *error* models an explicit abnormal termination.

Definition 5.1. Fix a global register set Reg and global space $H \cong \bigotimes_{x \in \text{Reg}} H_x$. *qwhile* commands are generated by the following grammar:

$$\begin{aligned} C \in \text{Cmd} ::= & \text{error} \mid \text{skip} \mid C_1; C_2 \mid \text{init } \rho_s \mid \text{apply } U_s \\ & \mid \text{if } (\Box m. M_s = m \rightarrow C_m) \text{ fi} \mid \text{while } M'_s = 1 \text{ do } C \text{ od.} \end{aligned}$$

Here: $s \subseteq \text{Reg}$ ranges over subsystems. ρ_s ranges over density operators on H_s (i.e. $\rho_s \succeq 0$, $\text{tr}(\rho_s) = 1$). U_s ranges over unitary operators on H_s . $M_s = \{(m, M_m)\}_{m \in \text{Out}(M_s)}$ is a measurement on H_s in Kraus form, i.e. $\sum_m M_m^\dagger M_m = \mathbf{1}_{H_s}$. $M'_s = \{M_0, M_1\}$ is a two-outcome measurement on H_s (a special case of the above).

skip terminates normally and leaves the state unchanged.

error terminates *abnormally*. It is intended to model bug signals such as failed runtime checks. Crucially, once *error* occurs, the program stops immediately: any remaining code is discarded rather than executed.

$C_1; C_2$ is sequential composition: execute C_1 first; if C_1 terminates normally then continue with C_2 ; if C_1 terminates abnormally then the whole composition terminates abnormally.

init ρ_s *resets* subsystem s to ρ_s , discarding any prior content of s (including entanglement with \bar{s}). Operationally, it applies $\rho \mapsto \rho_s \otimes \text{tr}_s(\rho)$.

apply U_s applies the unitary U_s to subsystem s (lifted cylindrically to H), i.e. $\rho \mapsto U_s^{(s)} \rho (U_s^{(s)})^\dagger$.

if($\Box m. M_s = m \rightarrow C_m$)*fi* first measures subsystem s using the measurement M_s . The (classical) outcome m selects the branch C_m , and the quantum state is updated by the corresponding Kraus operator M_m .

while $M'_s = 1$ *do* C *od* repeatedly measures subsystem s using $M'_s = \{M_0, M_1\}$. Outcome 0 terminates the loop; outcome 1 executes the body C and repeats.

5.2 Operational semantics

The semantics is given as a labelled transition system on configurations. Exit conditions are

$$\epsilon ::= \text{ok} \mid \text{er},$$

where **ok** denotes normal steps/termination and **er** denotes abnormal termination caused by error. We call outputs of **ok**-terminations *normal states* and outputs of **er**-terminations *abnormal states*.

A *configuration* is a pair $\langle C, \rho \rangle$ where C is the remaining code to be executed (or \downarrow to denote termination by convention) and $\rho \in \mathcal{D}^-(H)$ is the current program state. The one-step transition relation is written

$$\langle C, \rho \rangle \xrightarrow{\epsilon} \langle C', \rho' \rangle.$$

Because measurements branch on outcomes, a configuration may have multiple **ok**-successors, each carrying a different (subnormalized) post-measurement state.

Conventions. Whenever an operator is defined only on a subsystem s , it is understood to act on the whole space via cylindrical extension (e.g. $U_s^{(s)}$ and $M_m^{(s)}$). All rules below are to be understood under the fixed identification $H \cong H_s \otimes H_{\bar{s}}$ whenever subsystem s is involved, and with cylindrical extensions $A^{(s)} := A \otimes \mathbf{1}_{H_{\bar{s}}}$.

$$\frac{}{\langle \text{skip}, \rho \rangle \xrightarrow{\text{ok}} \langle \downarrow, \rho \rangle} (\text{SKIP}) \quad \frac{}{\langle \text{error}, \rho \rangle \xrightarrow{\text{er}} \langle \downarrow, \rho \rangle} (\text{ERROR})$$

The ERROR rule is the essence of abnormal termination: it stops execution immediately, raises label **er**, and returns the current quantum state unchanged.

$$\frac{\langle C_1; C_2, \rho \rangle \xrightarrow{\text{ok}} \langle C_2, \rho' \rangle}{\langle C_1, \rho \rangle \xrightarrow{\text{ok}} \langle \downarrow, \rho' \rangle} (\text{SEQ-DONE}) \quad \frac{\langle C_1; C_2, \rho \rangle \xrightarrow{\text{ok}} \langle C'_1; C_2, \rho' \rangle}{\langle C_1, \rho \rangle \xrightarrow{\text{ok}} \langle C'_1, \rho' \rangle} (\text{SEQ-STEP})$$

$$\frac{\langle C_1; C_2, \rho \rangle \xrightarrow{\text{er}} \langle \downarrow, \rho' \rangle}{\langle C_1, \rho \rangle \xrightarrow{\text{er}} \langle \downarrow, \rho' \rangle} (\text{SEQ-ERR})$$

These sequencing rules express two behaviors: (i) normal sequencing proceeds by stepping C_1 until it finishes, then continues with C_2 ; (ii) *er short-circuits* sequencing: if C_1 terminates abnormally, then the whole program terminates abnormally and C_2 is discarded.

$$\frac{}{\langle \text{init } \rho_s, \rho \rangle \xrightarrow{\text{ok}} \langle \downarrow, \rho_s \otimes \text{tr}_s(\rho) \rangle} (\text{INIT}) \quad \frac{}{\langle \text{apply } U_s, \rho \rangle \xrightarrow{\text{ok}} \langle \downarrow, U_s^{(s)} \rho (U_s^{(s)})^\dagger \rangle} (\text{APPLY})$$

The INIT rule makes explicit that initialization *forgets* whatever was stored in s by taking $\text{tr}_s(\rho)$ and then prepares ρ_s afresh. The APPLY rule is local unitary evolution: it preserves trace, so $\text{tr}(\rho)$ (the path weight) is unchanged.

$$\frac{\langle \text{if } (\square m. M_s = m \rightarrow C_m) \text{ fi}, \rho \rangle \xrightarrow{\text{ok}} \langle C_m, M_m^{(s)} \rho (M_m^{(s)})^\dagger \rangle}{m \in \text{Out}(M_s)} (\text{IF})$$

This rule formalizes *measurement-controlled branching*. For each possible outcome m , there is a transition to the corresponding branch C_m . The post-measurement state is subnormalized: its trace equals the probability mass of seeing outcome m on input ρ .

$$\frac{\langle \text{while } M'_s = 1 \text{ do } C \text{ od}, \rho \rangle \xrightarrow{\text{ok}} \langle \downarrow, M_0^{(s)} \rho (M_0^{(s)})^\dagger \rangle}{(\text{WHILE-0})}$$

$$\frac{\langle \text{while } M'_s = 1 \text{ do } C \text{ od}, \rho \rangle \xrightarrow{\text{ok}} \langle C; \text{ while } M'_s = 1 \text{ do } C \text{ od}, M_1^{(s)} \rho (M_1^{(s)})^\dagger \rangle}{(\text{WHILE-1})}$$

The loop is driven by a two-outcome measurement. Outcome 0 terminates the loop immediately; outcome 1 unrolls one iteration: execute C and then repeat the loop.

Comments on normalization and probability. The rules intentionally avoid renormalization. For example, $M_m^{(s)} \rho (M_m^{(s)})^\dagger$ is not divided by its trace. Instead, $\text{tr} \left(M_m^{(s)} \rho (M_m^{(s)})^\dagger \right)$ is the probability mass of taking outcome m from state ρ . This is why subnormalized states are convenient: the branching structure is classical (one successor per outcome), while the quantitative weights are carried by the traces of the resulting partial states.

Write $\xrightarrow{\epsilon}^*$ for the reflexive-transitive closure of $\xrightarrow{\epsilon}$, where the label records whether an abnormal termination occurred:

$\langle C, \rho \rangle \xrightarrow{\text{ok}}^* \langle C', \rho' \rangle$ means: execute zero or more steps, and no step is labelled **er**.

$\langle C, \rho \rangle \xrightarrow{\text{er}}^* \langle \downarrow, \rho' \rangle$ means: along the execution, an **er**-labelled termination occurs. By SEQ-ERR, once **er** happens, the remaining code is discarded and the computation ends immediately.

This separation of **ok** vs. **er** executions is essential for later incorrectness-style specifications: it allows us to talk about *normal outcomes* and *error outcomes* of the same program, and to treat the existence of an **er**-path as evidence of a bug.